


R data types: Lists

Alberto Garfagnini

Università di Padova

R lecture 3



R internals: variables and objects creation

- We create a vector with three values and assign it to a reference variable, `x`

```
x <- c(1,2,3)
```

- we now copy `x` to another variable `y`:

```
y <- x
```

- and modify one element of `y`

```
y[3] <- 4
```


- did we modify also `x`?

No, they refer to two different objects:

```
str(x)
%> num [1:3] 1 2 3
str(y)
%> num [1:3] 1 2 4
```

- the behavior is called **copy-on-modify**
- all R objects are **immutable**

```
lobstr::obj_addr(x)
"0x55d03cd66fb8"
```



```
lobstr::obj_addr(y)
"0x55d03dbac8c8"
```



The `lobstr` package allows to visualize R data structures: it shows memory location and size of objects.

URL: <https://github.com/r-lib/lobstr>

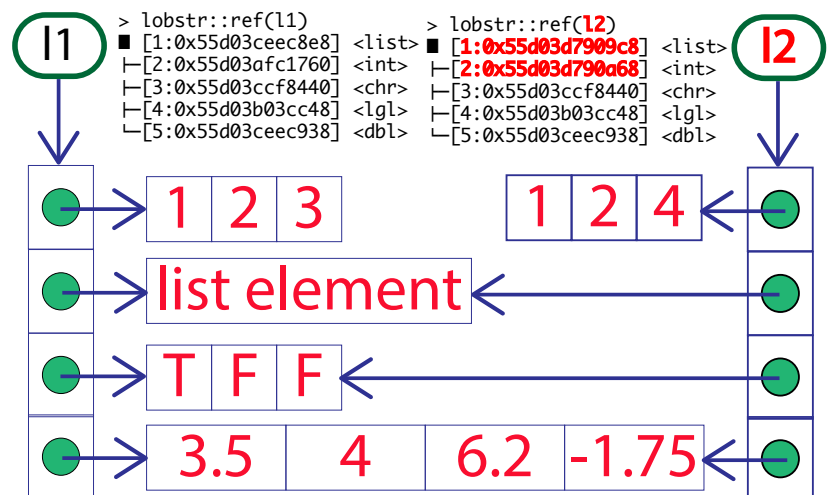
R Lists

- **Lists** are an evolution of atomic vectors: **each element can be of any type**
- from the technical point of view: **each element of a list** is of the same type: it is **a reference to another R object**
- building a list:

```
l1 <- list( 1:3,
            "list element",
            c(TRUE, FALSE, FALSE),
            c(3.5, 4, 6.2, -1.75)
)
typeof(l1)
%> [1] "list"
```

- we copy to a new list and modify one element

```
l2 <- l1
l2[[1]] <- c(1L, 2L, 4L)
```



R matrices

- a matrix is a 2-dimensional object
- the first way of creating a matrix is by calling the `matrix()` object constructor

```
X <- matrix(c(1,0,0,0,1,0,0,0,1), nrow=3) ; X
%>      [,1] [,2] [,3]
%> [1,]    1    0    0
%> [2,]    0    1    0
%> [3,]    0    0    1

class(X)
%> [1] "matrix"
attributes(X)
%> $dim
%> [1] 3 3
str(X)
%> num [1:3, 1:3] 1 0 0 0 1 0 0 0 1
```

- another way is to transform a vector in a matrix: data can be arranged by rows (`byrow=T`) or columns (`byrow=F`)

```
vct <- c(1,2,3,4,4,3,2,1)
V <- matrix(vct, byrow=T, nrow=2)
V
%>      [,1] [,2] [,3] [,4]
%> [1,]    1    2    3    4
%> [2,]    4    3    2    1
```

```
V <- matrix(vct, byrow=F, nrow=2)
V
%>      [,1] [,2] [,3] [,4]
%> [1,]    1    3    4    2
%> [2,]    2    4    3    1
```

- another possibility is to convert the vector to a matrix by specifying the new dimensions (rows and columns), using the `dim` function

```
vct <- c(1,2,3,4,4,3,2,1)
vct
%> [1] 1 2 3 4 4 3 2 1
```

```
dim(vct) <- c(4,2)
is.matrix(vct)
%> [1] TRUE
```

```
vct
%>      [,1] [,2]
%> [1,]    1    4
%> [2,]    2    3
%> [3,]    3    2
%> [4,]    4    1
```

- we can then transform the matrix:

```
tvct <- t(vct) # transpose the matrix
tvct
%>      [,1] [,2] [,3] [,4]
%> [1,]    1    2    3    4
%> [2,]    4    3    2    1
```

Accessing or operating on matrix rows or columns

- Let's create a matrix with $n = 20$ entries sampled from a Poisson distribution with $\lambda = 1.5$

```
X <- matrix(rpois(n=20,lambda=1.5), nrow=4)
X
%>      [,1] [,2] [,3] [,4] [,5]
%> [1,]    1    1    1    2    4
%> [2,]    1    1    3    3    2
%> [3,]    1    3    5    0    1
%> [4,]    2    1    1    2    2

X[3,3] # return element in row 3 and column 3
%> [1] 5
X[4,] # return row 4
%> [1] 2 1 1 2 2
X[,5] # return column 5
%> [1] 4 2 1 2
```

- there are special functions for calculating summary statistics on a matrix:

```
rowSums(X) # use colSums(X) for columns
%> [1] 9 10 10 8
rowMeans(X) # use colMeans(X) for columns
%> [1] 1.8 2.0 2.0 1.6
```

Adding rows and columns to a matrix

- given a matrix, we would like to add a row, at the bottom, showing the column means, and a column at the right showing the row variances:

```
vct <- matrix(c(1,0,2,5,1,1,3,1,3,1,0,2,1,0,2,1), byrow=T, nrow=4)
vct
%>      [,1] [,2] [,3] [,4]
%> [1,]    1    0    2    5
%> [2,]    1    1    3    1
%> [3,]    3    1    0    2
%> [4,]    1    0    2    1

vct <- rbind(vct, apply(vct, 2, mean))
vct <- cbind(vct, apply(vct, 1, var))

colnames(vct) <- c(1:4, "variance")
rownames(vct) <- c(1:4, "mean")

vct
%>      1    2    3    4  variance
%> 1    1.0 0.0 2.00 5.00 4.6666667
%> 2    1.0 1.0 3.00 1.00 1.0000000
%> 3    3.0 1.0 0.00 2.00 1.6666667
%> 4    1.0 0.0 2.00 1.00 0.6666667
%> mean 1.5 0.5 1.75 2.25 0.5416667
```

The apply() collection functions

- are used to apply operations on a the elements of a complex object (vector, list, data.frame, ...) avoiding the use of loops
- apply() is the most basic and can be used over a matrix or array

apply()

- usage:

```
      apply(X, MARGIN, FUN)

with

X: an array or matrix

MARGIN: a value or range between 1 and 2
        to define where to apply the function:
MARGIN=1: the manipulation is performed on rows
MARGIN=2: the manipulation is performed on columns
MARGIN=c(1,2) the manipulation is performed
              on rows and columns

FUN: which function to apply.
     Built functions like mean, median, sum, min, max
     and user-defined functions can be applied
```

The apply() collection functions

- `sapply()` takes a vector or list object and returns an object of the same type.

`sapply()`

- usage:

```
sapply(X, FUN)
```

Arguments:

X: A **vector** or an object

FUN: Function applied to each element of X

```
x <- 1:10
```

```
apply(x, 1, sqrt)
```

```
#> Error in apply(x, 1, sqrt) : dim(X) must have a positive length
```

```
sapply(x, sqrt)
```

```
[1] 1.000000 1.414214 1.732051 2.000000 2.236068
```

```
[6] 2.449490 2.645751 2.828427 3.000000 3.162278
```

The apply() collection functions

- `lapply()` performs operations on list objects and returns a list object with the same length of the original set
- `lapply()` applies a function to each element of the list
- `lapply()` takes list, vector or data frame as input and returns a list

`lapply()`

- usage:

```
lapply(X, FUN)
```

Arguments:

X: A **vector** or an object

FUN: Function applied to each element of X

```
box <- c("Orange", "CHerrry", "APPLE")
```

```
str(box)
```

```
#> chr [1:3] "Orange" "CHerrry" "APPLE"
```

```
lbox <- lapply(box, tolower)
```

```
str(lbox)
```

```
#> List of 3
```

```
#> $ : chr "orange"
```

```
#> $ : chr "cherrry"
```

```
#> $ : chr "apple"
```

apply(), sapply() and lapply()

- `apply()` is used to apply functions to rows or columns of matrices or dataframes across one of the margins of a matrix
- `margin=1` refers to the rows and `margin=2` to the columns

```
(Y <- matrix(rbinom(20, 9, 0.45), nrow=4))
%>      [,1] [,2] [,3] [,4] [,5]
%> [1,]    6    5    4    2    5
%> [2,]    6    3    3    5    4
%> [3,]    5    2    3    4    4
%> [4,]    3    4    4    3    6
apply(Y, MARGIN=2, FUN=sum) # apply sum() to all columns
%> [1] 20 14 14 14 19
```

- we can `apply()` functions to the individual elements of a matrix. In this case, the `margin` parameter, determines only the shape of the resulting matrix

```
apply(Y, 1, sqrt)
%>      [,1] [,2] [,3] [,4]
%> [1,] 2.449490 2.449490 2.236068 1.732051
%> [2,] 2.236068 1.732051 1.414214 2.000000
%> [3,] 2.000000 1.732051 1.732051 2.000000
%> [4,] 1.414214 2.236068 2.000000 1.732051
%> [5,] 2.236068 2.000000 2.000000 2.449490
apply(Y, 2, sqrt)
%>      [,1] [,2] [,3] [,4] [,5]
%> [1,] 2.449490 2.236068 2.000000 1.414214 2.236068
%> [2,] 2.449490 1.732051 1.732051 2.236068 2.000000
%> [3,] 2.236068 1.414214 1.732051 2.000000 2.000000
%> [4,] 1.732051 2.000000 2.000000 1.732051 2.449490
```

apply(), sapply() and lapply()

- it is also possible to apply an anonymous, user defined, function

```
apply(Y, 1, function(x) x^2+x) # compute x^2 + x for each element
%>      [,1] [,2] [,3] [,4]
%> [1,]   42   42   30   12
%> [2,]   30   12    6   20
%> [3,]   20   12   12   20
%> [4,]    6   30   20   12
%> [5,]   30   20   20   42
```

- in case you need to apply a function to a vector, rather than to the margin of a matrix, use `sapply()`

```
sapply(12:14, seq) # generate a list of seq, from 1:12 to 1:14
%> [[1]]
%> [1]  1  2  3  4  5  6  7  8  9 10 11 12
%>
%> [[2]]
%> [1]  1  2  3  4  5  6  7  8  9 10 11 12 13
%>
%> [[3]]
%> [1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14
```

- random numbers from a uniform distribution $\mathcal{U}(0, 1)$ are generated using `runif()`
- the random generation seed is set via `set.seed()`

```
set.seed(2019)
runif(3)
%> [1] 0.7699015 0.7128397 0.3033602
```

- resetting the seed with the same value will generate the same sequence of random numbers
- it is also possible to save the current seed and reuse it to obtain the same random numbers sub-sequence

```
current.seed <- .Random.seed # save the current seed
runif(3)
%> [1] 0.61823636 0.05048374 0.04321880
runif(3)
%> [1] 0.820176206 0.009614496 0.102491504

current.seed -> .Random.seed # reset the previous sequence seed
runif(5)
%> [1] 0.618236361 0.050483740 0.043218804 0.820176206 0.009614496
```

sampling from a vector

- generating random numbers from probability distributions will be discussed in the next lessons
- now we want to randomize (shuffling or sampling from) the elements of a vector
- There are two ways of sampling:
 - 1) **sampling without replacement** : all the vector values will appear in output, but in a randomized sequence
 - 2) **sampling with replacement** : some vector values may be re-selected and appear more than once in the output
- using `sample()`, sampling without replacement is the default operation

```
y <- c(8,3,5,7,6,6,8,9,2,3,9,4,10,4,11)
sample(y) # reshuffling all vector values
%> [1] 3 9 2 5 4 8 6 8 6 4 10 3 7 11 9
sample(y, 5) # pick up only 5 values from the original vector
%> [1] 3 8 9 4 8
sample(y, 5) # just redo it, and a different sequence may appear
%> [1] 8 3 8 4 3
```

- The option `replace=T` allows for sampling with replacement

```
sample(y, replace=T)
%> [1] 8 3 6 8 8 4 3 7 10 9 10 9 4 4 7
```

sample()'s surprise example

```
x <- 1:10
x
%> [1] 1 2 3 4 5 6 7 8 9 10
```

```
sample(x[x>8])
%> [1] 10 9
```

sample(x, size, replace = FALSE, prob = NULL)
If 'x' has length 1, sampling takes place from '1:x'

```
sample(x[x>9])
%> [1] 1 10 8 7 6 5 4 2 9 3
```

```
sample(x[x > 10])
%> integer(0)
```

- the first argument of `sample()` can be a vector of more than one element or an integer
- the `resample()` function is safer

```
sample(15)
%> [1] 2 4 3 11 7 14 6 5 1 13 15 10 12 9 8
```

```
str(x[x>9])
%> int 10
```

the `resample()` function is available in the `gdata` package. <https://cran.r-project.org/web/packages/gdata/index.html>

```
resample(x[x>8])
%> [1] 10 9
resample(x[x>9])
%> [1] 10
```

R subsetting

- R's subsetting operators are fast and powerful, and allow to perform complex operations in a way that few other languages can match
 - there are 6 ways to subset atomic vectors
 - there are 3 subsetting operators: `[[`, `[` and `$`
 - subsetting can be combined with assignment

Subsetting atomic vectors - 1

```
x <- c(2.1, 4, 6.7, 1.75)
```

- **positive integers** return elements at a specified position

```
x[c(1,3)]  
%> [1] 2.1 6.7
```

```
% Duplicate indices will duplicate values  
x[c(1,1,3,3)]  
%> [1] 2.1 2.1 6.7 6.7
```

```
% Real numbers are truncated to integers  
x[sort(x)]  
%> [1] 2.10 4.00 1.75 NA
```

- **negative integers** exclude elements

```
x[-c(1,3)]  
%> [1] 4.00 1.75
```

```
% NB negative and positive ints cannot be mixed  
x[c(-1,3)]  
%> Error in x[c(-1, 3)]: only 0's may be mixed with negative subscripts
```

Subsetting atomic vectors - 2

```
x <- c(2.1, 4, 6.7, 1.75)
```

- **logical vectors** select elements where the logical value is **TRUE**

```
x[c(T, T, F, T)]  
%> [1] 2.10 4.00 1.75
```

```
x[x>2]  
%> [1] 2.1 4.0 6.7
```

- if in `x[sel]`, `length(sel) != length(x)` the **recycling rules** are used: the shorter vector is recycled to the length of the longer

```
> x[c(TRUE, FALSE)]  
[1] 2.1 6.7
```

```
%# is equivalent to:  
> x[c(TRUE, FALSE, TRUE, FALSE)]  
[1] 2.1 6.7
```

- **nothing** returns the original vector

```
x[]  
%> [1] 2.10 4.00 6.70 1.75
```

Subsetting atomic vectors - 3

```
x <- c(2.1, 4, 6.7, 1.75)
```

- `zero` returns a zero-length vector (it can be helpful to generate test data)

```
x[0]  
numeric(0)
```

- `named vectors` can be accessed with `character vectors`

```
y <- setNames(x, LETTERS[1:length(x)])  
y  
%>      A      B      C      D  
%> 2.10 4.00 6.70 1.75  
y["A"]  
%>      A  
%> 2.1  
  
y[c('A', 'A', 'D')]  
%>      A      A      D  
%> 2.10 2.10 1.75
```

- WARNING: subsetting with factors will use the underlying integer vector, not the character levels. → [Avoid subsetting with factors](#)

```
y[factor("B")]  
%>      A  
%> 2.1
```

Subsetting matrices

- subsetting a matrix or a list works in a similar way as subsetting atomic vectors

```
S <- matrix(1:9, nrow = 3)  
%> [1,] 1 4 7  
%> [2,] 2 5 8  
%> [3,] 3 6 9
```

- using `[]` always `returns a list`
- `[[]` and `$` allows to `pull out elements from the list`
- the common rule to subset a matrix (2D) and an array (nD , $n > 2$) is to supply a 1D vector for each dimension, separated by a comma
- `blank subsetting` allows to keep all data for the corresponding dimension

```
%# Get rows 1 and 3 and all columns  
S[c(1,3), ]  
%>      [,1] [,2] [,3]  
%> [1,] 1 4 7  
%> [2,] 3 6 9  
  
colnames(S) <- c("S1", "S2", "S3")  
S[c(T, F, T), c("S1", "S3")]  
%>      S1 S3  
%> [1,] 1 7  
%> [2,] 3 9
```

Subsetting matrices - 2

- matrices and arrays are just vectors with special attributes, therefore they can be subset with a single vector, as if they were a 1D vector

```
v <- outer(1:5, 1:5, FUN="paste", sep=",")
v
%>      [,1] [,2] [,3] [,4] [,5]
%> [1,] "1,1" "1,2" "1,3" "1,4" "1,5"
%> [2,] "2,1" "2,2" "2,3" "2,4" "2,5"
%> [3,] "3,1" "3,2" "3,3" "3,4" "3,5"
%> [4,] "4,1" "4,2" "4,3" "4,4" "4,5"
%> [5,] "5,1" "5,2" "5,3" "5,4" "5,5"

v[seq(3, 23, 5)]
%> [1] "3,1" "3,2" "3,3" "3,4" "3,5"
```

- to preserve the original matrix dimension, use `drop = FALSE`

```
(S <- matrix(1:6, nrow = 2))
%>      [,1] [,2] [,3]
%> [1,]    1    3    5
%> [2,]    2    4    6

S[1, ]
%> [1] 1 3 5

S[1, , drop = FALSE]
%>      [,1] [,2] [,3]
%> [1,]    1    3    5
```

Selecting a single element


- there are two other subsetting operators:
 - `[[]]` is used to extract single items
 - `$` is used as a shorthand: `x$y` stands for `x[["y"]]`
- `[[]]` is most important while working with lists: subsetting a list with single `[[]]` always returns a smaller list


If list `xl` is a train carrying objects, then `xl[[5]]` is the object in car 5; `xl[4:6]` is a train of cars 4-6

<https://twitter.com/RLangTip/status/268375867468681216>


- with this metaphor let's build a list

```
xl <- list(1:3, "one", c(T,F,F))
```









- two options are available when extracting a single element:
 - create a smaller train, with fewer cars (using `[]`)



- or extract the content of a particular car (with `[[]]`)



- extracting multiple (or zero) elements, we have to build a smaller train



Loops

- Let's create a function, using loops, to evaluate the factorial:
- $n! = n \cdot (n - 1) \cdot (n - 2) \dots 2 \cdot 1$

```
fac1 <- function(x) {
  f <- 1
  if (x<2) return (1)
  for (i in 2:x) {
    f <- f*i
  }
  return (f)
}
sapply(1:5, fac1)
%> [1] 1 2 6 24 120
```

```
fac2 <- function(x) {
  f <- 1; t <- x
  while (t>1) {
    f <- f*t
    t <- t-1
  }
  return(f)
}
sapply(1:5, fac2)
%> [1] 1 2 6 24 120
```

```
fac3 <- function(x) {
  f <- 1; t <- x
  repeat {
    if (t<2) break
    f <- f*t
    t <- t-1
  }
  return(f)
}
sapply(1:5, fac3)
%> [1] 1 2 6 24 120
```

- But it is almost always better to use a built-in function that operates on the entire vector, removing the need of loops or repeats

```
> cumprod(1:5) # it does not work for 0
[1] 1 2 6 24 120
> fac4 <- function(x) max(cumprod(1:x))
> sapply(1:5, fac4)
[1] 1 2 6 24 120
```

```
> # R implements a factorial() function, introduced not long ago
> sapply(1:5, factorial)
[1] 1 2 6 24 120
```

- it's a good R programming practice to avoid loops wherever possible
- in many cases, using vector functions, makes it particularly straightforward

```
> y <- c(-3,4,-2,-1,8,7,9)
> y
[1] -3  4 -2 -1  8  7  9

> for (i in 1:length(y)) {if (y[i] < 0) y[i] <- 0}
> y
[1] 0 4 0 0 8 7 9
```

- in the example below, a loop can be replaced by logical subscripts

```
> y <- c(-3,4,-2,-1,8,7,9)

> y[y<0] <- 0

> y
[1] 0 4 0 0 8 7 9
```

the ifelse() vectorized function

- ifelse() allow to work on an entire vector without using loops

```
> y <- log(rpois(20,1.5))
> y
[1]      -Inf 1.0986123 0.0000000 0.0000000 0.0000000 0.6931472
[7] 0.6931472      -Inf 1.3862944 0.6931472 1.3862944      -Inf
[13]      -Inf 0.0000000 0.0000000 1.0986123 0.0000000 0.0000000
[19]      -Inf 1.0986123

> mean(y)
[1] -Inf

> (y <- ifelse(y<0, NA, y))
[1]      NA 1.0986123 0.0000000 0.0000000 0.0000000 0.6931472
[7] 0.6931472      NA 1.3862944 0.6931472 1.3862944      NA
[13]      NA 0.0000000 0.0000000 1.0986123 0.0000000 0.0000000
[19]      NA 1.0986123

> mean(y, na.rm=TRUE)
[1] 0.5431911
```

Loops are slow, compared to vectorized operations

- let's generate $5 \cdot 10^7$ events according to an uniform distribution, $\mathcal{U}(0,1)$
- we want to search for the maximum value in the vector using the vectorized function `max()` or by using conventional loops

```
x <- runif(50000000)

system.time(max(x))
%>   user   system elapsed
%> 0.106   0.000   0.106

pc <- proc.time()
cmax <- x[1]
for (i in 2:length(x)) { if(x[i]>cmax) cmax <- x[i] }

proc.time()-pc
%>   user   system elapsed
%> 2.061   0.071   2.133
```

- `system.time()` and `proc.time()` produce a vector of three numbers, showing the user, system and total elapsed time in seconds

Good/Bad practice in building vectors

- we want to build a vector containing 10^n elements in the sequence $1:10^n$
- three ways are analyzed


```
test1 <- function(n){
  y <- 1:n
}

test2 <- function(n){
  y <- numeric(n)
  for (i in 1:n)
    y[i] <- i
}

test3 <- function(n){
  y <- NULL
  for (i in 1:n)
    y <- c(y,i)
}
```

```
> system.time(test1(10000000))
  user   system elapsed
0.006   0.000   0.006
> system.time(test2(10000000))
  user   system elapsed
0.622   0.011   0.633
> system.time(test3(10000000))
  user   system elapsed
2.755   0.003   2.758
```

- the first method (test1) is the best
- the loop using a pre-determined vector length is reasonably fast
- the last method (test3) is the slowest
- Moral: **never grow vectors by repeated concatenation**



- let's create a more complex list with heterogeneous object types

```
apples  <- c(4, 4.5, 5, 3.9)
oranges <- c(TRUE, FALSE, TRUE)
chalk   <- c("limestone", "marl", "ooline", "CaCO3")
pears   <- c(3.2-4.5i, 12.8+2.2i)

items <- list(apples, oranges, chalk, pears)
items
%> [[1]]
%> [1] 4.0 4.5 5.0 3.9
%>
%> [[2]]
%> [1] TRUE FALSE TRUE
%>
%> [[3]]
%> [1] "limestone" "marl"          "ooline"        "CaCO3"
%>
%> [[4]]
%> [1] 3.2-4.5i 12.8+2.2i
```

R List example: element access

- vectors, matrices and arrays subscripts have one set of square brackets [6], [3,4] or [2,3,2,1]
- lists subscripts have double square brackets [[2]] or [[i,j]]

```
items[[3]]
%> [1] "limestone" "marl"          "ooline"        "CaCO3"

items[[3]][1]
%> [1] "limestone"
```

- if the list elements have names, it is possible to use the operator \$ for list indexing

```
names(items) <- c("apples", "oranges", "chalk", "pears")

str(items)
%> List of 4
%> $ apples : num [1:4] 4 4.5 5 3.9
%> $ oranges: logi [1:3] TRUE FALSE TRUE
%> $ chalk  : chr [1:4] "limestone" "marl" "ooline" "CaCO3"
%> $ pears  : cplx [1:2] 3.2-4.5i 12.8+2.2i

items$pears
%> [1] 3.2-4.5i 12.8+2.2i
```

- the length of the list is the number of items on the list. To get the length of the individual vectors we use the `length()` function

```
length(items)
%> [1] 4

lapply(items, length)
%> $apples
%> [1] 4

%> $oranges
%> [1] 3

%> $chalk
%> [1] 4

%> $pears
%> [1] 2

class(items)
%> [1] "list"

lapply(items, class)
%> $apples
%> [1] "numeric"

%> $oranges
%> [1] "logical"

%> $chalk
%> [1] "character"

%> $pears
%> [1] "complex"
```

R Lists : 3

- applying numeric functions to the list, will only work for objects of class `numeric` or `complex`

```
mean(items)
%> [1] NA
%> Warning message:
%> In mean.default(items) :
%>   argument is not numeric or logical: returning NA

lapply(items, mean)
%> $apples
%> [1] 4.35

%> $oranges
%> [1] 0.6666667

%> $chalk
%> [1] NA

%> $pears
%> [1] 8-1.15i
%> Warning message:
%> In mean.default(X[[i]], ...) :
%>   argument is not numeric or logical: returning NA
```

- a warning message points out that the third vector cannot be coerced to a number (it is not numeric, complex or logical), therefore NA appears in the output

- The `summary()` function works for lists, but the most useful overview of a list content is given by `str()`, the structure function:

```
summary(items)
%>      Length Class  Mode
%> apples    4      -none- numeric
%> oranges    3      -none- logical
%> chalk      4      -none- character
%> <NA>       2      -none- complex

str(items)
%> List of 4
%> $ apples : num [1:4] 4 4.5 5 3.9
%> $ oranges: logi [1:3] TRUE FALSE TRUE
%> $ chalk  : chr [1:4] "limestone" "marl" "ooline" "CaCO3"
%> $ NA     : cplx [1:2] 3.2-4.5i 12.8+2.2i
```

Question Time

- What is the effect of `[[1:2]]` on a list ?

```
x1 <- list(1:3,
           "one",
           c(T, F, F))

% x1[[1:2]] is equivalent to x1[[1]][[2]]
x1[[1:2]]
[1] 2
x1[[1]][[2]]
%> [1] 2
% i.e. it extracts the element stored in the list at position 1, and
%      then it gets the second element

% This fails:
x1[[1:3]]
%> Error in x1[[1:3]] : recursive indexing failed at level 2
% and it is equivalent to:
> (a1 <- x1[[1]])
[1] 1 2 3
> (a2 <- a1[[2]])
[1] 2
> (a3 <- a2[[3]])
Error in a2[[3]] : subscript out of bounds
```