


R data types: Data Frames

Alberto Garfagnini

Università di Padova

R lecture 4



R Data frames

- two important S3 vectors built on top of lists are **data frames** and **tibbles**
- a data frame is like a matrix, with a 2-dim rows-and-columns structure
- it's a **named list of vectors**, with attributes for columns and rows names, (**names**, **row.names**), belonging to the **data.frame** class
- **technically, a data frame is a list with all equal length vectors**


```
df1 <- data.frame(x = 1:3, y = letters[1:3])
typeof(df1)
%> [1] "list"

attributes(df1)
%> $names
%> [1] "x" "y"

%> $class
%> [1] "data.frame"

%> $row.names
%> [1] 1 2 3

str(df1)
%> 'data.frame':      3 obs. of  2 variables:
%> $ x: int  1 2 3
%> $ y: Factor w/ 3 levels "a","b","c": 1 2 3
```



R Data Frames : examples

- we have a table with the results of two exams for the student of an hypothetical course, and we want to import them in a `data.frame`

Exam1	Exam2	Gender
27	25	M
28	30	F
...		
27	27	M
25	28	F

```
exam1 <- c(27,28,24,24,30,26,23,23,24,28,27,25)
exam2 <- c(25,30,26,24,30,30,25,25,30,28,27,28)
gender <- c("M","F","M","M","M","M","M","M","F","F","M","F")

dc <- data.frame(exam1, exam2, gender)
head(dc, n=2) # extract the first two lines of the data frame
%>   exam1 exam2 gender
%> 1    27    25      M
%> 2    28    30      F
```

From R 4.0
stringsAsFactors = FALSE
by default

```
dc1 <- data.frame(exam1, exam2, gender,
                  stringsAsFactors = FALSE)
str(dc1)
'data.frame':   12 obs. of  3 variables:
 $ exam1 : num  27 28 24 24 30 26 23 23 24 28 ...
 $ exam2 : num  25 30 26 24 30 30 25 25 30 28 ...
 $ gender: chr  "M" "F" "M" "M" ...
```

R Data Frames objects creation

- Data frames are list of vectors, therefore copy-on-modify has important consequences


```
d1 <- data.frame(x = c(1, 2, 3),
                 y = c(5, 7, 9))

d1
%>   x y
%> 1 1 5
%> 2 2 7
%> 3 3 9
```


- if we modify a column → only the reference to the new column will be updated

```
d2 <- d1
d2[, 2] <- d2[, 2] * 2
d2
%>   x y
%> 1 1 10
%> 2 2 14
%> 3 3 18
```

```
> lobstr::ref(d1)
[1:0x55905d24e9e8] <df[,2]>
└─x = [2:0x55905e564eb8] <dbl>
└─y = [3:0x55905e564eb8] <dbl>
```



```
> lobstr::ref(d1)
[1:0x55905d24e9e8] <df[,2]>
└─x = [2:0x55905e564eb8] <dbl>
└─y = [3:0x55905e564eb8] <dbl>
```



```
> lobstr::ref(d2)
[1:0x55905c9f1628] <df[,2]>
└─x = [2:0x55905e564eb8] <dbl>
└─y = [3:0x55905e92cd78] <dbl>
```


R Data Frames objects creation

- Data frames are list of vectors, therefore copy-on-modify has important consequences

```
d1 <- data.frame(x = c(1, 2, 3),  
                 y = c(5, 7, 9))
```

```
d1  
%>   x y  
%> 1 1 5  
%> 2 2 7  
%> 3 3 9
```


```
> lobstr::ref(d1)  
■ [1:0x55905d24e9e8] <df[,2]>  
└─x = [2:0x55905e564eb8] <dbl>  
└─y = [3:0x55905e564e68] <dbl>
```



- but if any row is modified → every column is modified because every column must be copied

```
d3 <- d1  
d3[1, ] <- d3[1, ] * 3  
d3  
%>   x y  
%> 1 3 15  
%> 2 2 7  
%> 3 3 9
```


```
> lobstr::ref(d3)  
■ [1:0x55905e6f1058] <df[,2]>  
└─x = [2:0x55905ea0c238] <dbl>  
└─y = [3:0x55905ea0c1e8] <dbl>
```




Modify-in-place

- Modifying an R object usually **creates a copy**
- but there are **2 exceptions**:
 - objects with single binding** get a special performance optimization
 - environments**, a special type of object, **are always modified in place**

```
v <- c(1, 3, 2)  
lobstr::obj_addr(v)  
%> [1] "0x55905ea0f4e8"
```



```
v[[3]] <- -2  
lobstr::obj_addr(v)  
%> [1] "0x55905ea0f4e8"
```



- but **it is very difficult to predict** when R applies this **optimization**
- concerning object binding, **R only counts 0, 1 or MANY**
- it means that if an object has 2 bindings (i.e. many), and one gets deleted, the reference does not go back to 1 (many - 1 = many)
- when a function is called, it makes a reference to the object → it is very difficult to predict whether or not a copy will occur
- cfr: <https://developer.r-project.org/Refcnt.html>

Accessing data frames elements

- a data frame is a list, therefore we can access them via **component index** value `[[j]]` or via **component names**

```
str(dc)
%> 'data.frame': 12 obs. of 3 variables:
%> $ exam1 : num 27 28 24 24 30 26 23 23 24 28 ...
%> $ exam2 : num 25 30 26 24 30 30 25 25 30 28 ...
%> $ gender: Factor w/ 2 levels "F","M": 2 1 2 2 2 2 2 2 1 1 ...

dc[[1]] # access by component index
%> [1] 27 28 24 24 30 26 23 23 24 28 27 25

dc$exam1 # access by component name
%> [1] 27 28 24 24 30 26 23 23 24 28 27 25
%> Levels: F M
```

- but a data frame can be treated in a matrix-like fashion, as well

```
dc[,1] # select column 1
%> [1] 27 28 24 24 30 26 23 23 24 28 27 25

dc[1,1] # and access the single element, as well
%> [1] 27
```

Data frames row names

- data frames allow to **label each row with a name**, a character vector containing only unique names

```
df1 <- data.frame( age = c(35, 25, 18),
                   hair = c("blond", "brown", NA),
                   row.names = c("Bob", "Tom", "Sam"))

df1
%>      age  hair
%> Bob   35 blond
%> Tom   25 brown
%> Sam   18  <NA>

names(df1)
%> [1] "age" "hair"

row.names(df1)
%> [1] "Bob" "Tom" "Sam"
```

- but row names are a bad practice:

- (1) **metadata is metadata** : storing it in a different way to the rest of data is a bad idea
- (2) **row names are a poor abstraction for labeling rows** : they only work when a row can be identified by a single string
- (3) **row names must be unique** : any duplication of rows will create new row names
→ complicated "string surgery" may be needed

```
dc[2:4,] # Select only rows 2:4
%>      exam1 exam2 gender
%> 2      28     30      F
%> 3      24     26      M
%> 4      24     24      M
```

```
dc[-(2:10),] # drop rows 2:10
%>      exam1 exam2 gender
%> 1      27     25      M
%> 11     27     27      M
%> 12     25     28      F
```

- with the `sample` function , data can be selected at random

```
dc[sample(1:12,3),] # select 3 rows at random
%>      exam1 exam2 gender
%> 8      23     25      M
%> 9      24     30      F
%> 6      26     30      M
```

```
dc[sample(1:12,3),] # select 3 rows at random
%>      exam1 exam2 gender
%> 1      27     25      M
%> 10     28     28      F
%> 2      28     30      F
```

Advanced data frames : data selection

- suppose we want to extract all columns that contain numbers, rather than characters or logicals, from a data frame

```
dc[,sapply(dc,is.numeric)]
%>      exam1 exam2
%> 1      27     25
%> 2      28     30
%> 3      24     26
%> 4      24     24
%> 5      30     30
%> 6      26     30
%> 7      23     25
%> 8      23     25
%> 9      24     30
%> 10     28     28
%> 11     27     27
%> 12     25     28
```

```
dc <- data.frame(exam1, exam2, gender)
str(dc)
'data.frame':  12 obs. of  3 variables:
 $ exam1 : num  27 28 ...
 $ exam2 : num  25 30 ...
 $ gender: Fact w/ 2 levels "F","M": 2 1
```

- and now we want to get only factors (and remove numerics)

```
dc[,sapply(dc,is.factor)]
%> [1] M F M M M M M F F M F
%> Levels: F M
```

Advanced data frames and NA elements

- sometimes our data frame can have missing values (NA) and we may need to omit those values
- we can create a [shorter data frame](#) using the `na.omit()` function

```
data
%>   slope  pH area
%> 1    11 4.1  3.6
%> 2     NA 5.2  5.1
%> 3     3 4.9  2.8
%> 4     5  NA  3.7
```

```
na.omit(data)
%>   slope  pH area
%> 1    11 4.1  3.6
%> 3     3 4.9  2.8
```

```
clean_data <- na.exclude(data)
clean_data
%>   slope  pH area
%> 1    11 4.1  3.6
%> 3     3 4.9  2.8
```

```
lapply(clean_data, mean)
%> $slope
%> [1] 7
%> $pH
%> [1] 4.5
%> $area
%> [1] 3.2
```

```
# Let's count the missing values
apply(apply(data, 2, is.na), 2, sum)
%> slope    pH    area
%>      1      1      0
```

Advanced data frames : sorting elements

```
dc[order(exam1),]
%>   exam1 exam2 gender
%> 7     23    25      M
%> 8     23    25      M
...

```

```
dc[order(exam1, decreasing=TRUE),]
%>   exam1 exam2 gender
%> 5     30    30      M
%> 2     28    30      F

```

- `dc[order(gender, exam1, exam2, decreasing=TRUE),]`
%> exam1 exam2 gender
%> 5 30 30 M
%> 11 27 27 M
%> 1 27 25 M
%> 6 26 30 M
%> 3 24 26 M
%> 4 24 24 M
%> 7 23 25 M
%> 8 23 25 M
%> 2 28 30 F
%> 10 28 28 F
%> 12 25 28 F
%> 9 24 30 F

Summary of data selection in data frames

- given a data frame called `data`, we assume `n` is a row number, and `m` is one of the column.
- the syntax `[n,]` selects all the columns given row `n`, while `[,m]` selects all the rows with column `m`

command	meaning
<code>data[n,]</code>	select all of the columns from row <code>n</code> of the data frame
<code>data[-n,]</code>	drop the whole of row <code>n</code> from the data frame
<code>data[1:n,]</code>	select all of the columns from rows 1 to <code>n</code> of the data frame
<code>data[-(1:n),]</code>	drop all of the columns from rows 1 to <code>n</code> of the data frame
<code>data[c(i,j,k),]</code>	select all of the columns from rows <code>i</code> , <code>j</code> , and <code>k</code> of the data frame
<code>data[x > y,]</code>	use a logical test (<code>x > y</code>) to select all columns from certain rows
<code>data[,m]</code>	select all of the rows from column <code>m</code> of the data frame
<code>data[, -m]</code>	drop the whole of column <code>m</code> from the data frame
<code>data[, 1:m]</code>	select all of the rows from columns 1 to <code>m</code> of the data frame
<code>data[, -(1:m)]</code>	drop all of the rows from columns 1 to <code>m</code> of the data frame
<code>data[, c(i,j,k)]</code>	select all of the rows from columns <code>i</code> , <code>j</code> , and <code>k</code> of the data frame
<code>data[, x > y]</code>	use a logical test (<code>x > y</code>) to select all rows from certain columns
<code>data[, c(1:m,i,j,k)]</code>	add duplicate copies of columns <code>i</code> , <code>j</code> , and <code>k</code> to the data frame
<code>data[x > y, a != b]</code>	extract certain rows (<code>x > y</code>) and certain columns (<code>a != b</code>)
<code>data[c(1:n,i,j,k),]</code>	add duplicate copies of rows <code>i</code> , <code>j</code> , and <code>k</code> to the data frame

The tibble data structure

- it is a modern reimagining of the data frame
- it is provided by the tibble package which is part of the [tidyverse core library](https://tidyverse.org/)

```
library(tidyverse)
%> Attaching packages: tidyverse 1.3.0
%>   ggplot2 3.3.0    purrr   0.3.3
%>   tibble  2.1.3    dplyr   0.8.5
%>   tidyr   1.0.2    stringr 1.4.0
%>   readr   1.3.1    forcats 0.5.0
%> Conflicts: tidyverse_conflicts()
%> dplyr::filter() masks stats::filter()
%> dplyr::lag()    masks stats::lag()
```

- a data frame can be converted to a tibble

```
dct <- tibble(dc)
dct
%> # A tibble: 12 x 1
%>   dc$exam1 $exam2 $gender
%>   <dbl>   <dbl> <fct>
%> 1      27     25 M
%> ...
%> 12     25     28 F
```



<https://tibble.tidyverse.org/>

- or created from vectors (as for the data frame)

```
dct <- data.frame(exam1 = c(27,28,24,24,30,26,23,23,24,28,27,25),
                  exam2 = exam2, gender)
```

Tibbles vs. data.frame : printing

- two main differences in the usage of a tibble versus a data.frame:
- printing and subsetting
- Tibbles have a refined print method that shows only the first 10 rows:

```
atb <- tibble( a = lubridate::now() + runif(1e3) * 86400,
              b = 1:1e3,
              c = runif(1e3),
              d = sample(letters, 1e3, replace = TRUE) )

atb
%> # A tibble: 1,000 x 4
%>       a                b          c d
%>   <dtm>             <int>   <dbl> <chr>
%> 1 2020-03-24 08:56:26     1  1.00   q
%> 2 2020-03-24 10:57:29     2  0.996  z
%> 3 2020-03-24 00:32:33     3  0.620  r
%> 4 2020-03-24 01:16:23     4  0.804  d
%> 5 2020-03-24 03:32:17     5  0.311  e
%> 6 2020-03-24 00:22:27     6  0.206  u
%> 7 2020-03-23 12:58:48     7  0.0390 e
%> 8 2020-03-23 20:36:03     8  0.449  d
%> 9 2020-03-24 01:29:00     9  0.271  r
%> 10 2020-03-24 10:52:01    10  0.460  v
%> # with 990 more rows
```

Tibbles vs. data.frame : subsetting

- tibbles can extract by name or position

```
tb1 <- tibble( x = runif(5),
              y = rnorm(5) )

# Extract by name
tb1$x
#> [1] 0.7330 0.2344 0.6604 0.0329 0.4605
tb1[["x"]]
#> [1] 0.7330 0.2344 0.6604 0.0329 0.4605

# Extract by position
tb1[[1]]
#> [1] 0.7330 0.2344 0.6604 0.0329 0.4605
```

- tibbles are more strict than data.frame: they never do partial matching, and they will generate a warning if the column you are trying to access does not exist

```
tdc <- as_tibble(dc)
```

```
dc$gender
#> [1] M F M ... F M F
#> Levels: F M
```

```
dc$gen
#> [1] M F M ... F M F
#> Levels: F M
```

```
tdc$gender
#> [1] M F M ... F M F
#> Levels: F M
```

```
tdc$gen
#> NULL
#> Warning message:
#> Unknown or uninitialised column:
#> 'gen'.
```

Data Input

Data Input

- numbers can be [input](#)ed through the [keyboard](#), from the [Clipboard](#), from an external [file on disk](#), or from an external [file on the Web](#)
- use the concatenate function for up to 10 numbers
- and [scan\(\)](#) for typing or pasting data into a vector

```
y <- c (6,7,3,4,8,5,6,2)
```

```
tu <- scan()  
%> 1: 6  
%> 2: 3  
%> 3: 4  
%> 4: 2  
%> 5:  
%> Read 4 items  
tu  
%> [1] 6 3 4 2
```

- but the easiest way is to [read data from a file](#) (or from the Web), already shaped in a data frame format

Data Input using `read.table()`

- the `read.table()` function reads data from a local file and creates a data frame

```
data <- read.table("yield.txt",header=T)
```

```
data
%>      year wheat barley oats rye corn
%> 1  1980   5.9   4.4  4.1 3.8  4.4
%> 2  1981   5.8   4.4  4.3 3.7  4.1
%> 3  1982   6.2   4.9  4.4 4.1  4.0
...
%> 27 2006   8.0   5.9  6.0 6.1  4.5
%> 28 2007   7.2   5.7  5.5 5.7  3.9
%> 29 2008   8.3   6.0  5.8 6.1  4.4
```

- the parameter `header = T` tells R to use the first row as column names

```
names(data)
%> [1] "year"    "wheat"   "barley"  "oats"    "rye"     "corn"

str(data)
%> 'data.frame':      29 obs. of  6 variables:
%>  $ year   : int  1980 1981 1982 1983 1984 1985 1986 1987 1988 1989 ...
%>  $ wheat  : num  5.9 5.8 6.2 6.4 7.7 6.3 7 6 6.2 6.7 ...
%>  $ barley : num  4.4 4.4 4.9 4.7 5.6 5 5.2 5 4.7 4.9 ...
%>  $ oats   : num  4.1 4.3 4.4 4.3 4.9 4.6 5.2 4.6 4.6 4.5 ...
%>  $ rye    : num  3.8 3.7 4.1 3.7 4.7 4.6 4.7 4.8 4.6 4.8 ...
%>  $ corn   : num  4.4 4.1 4 4.1 4.7 4.3 4.3 4.5 4.2 3.8 ...
```

Data Input using `read.table()`

- if the separator between variable names and data fields are **not blanks** or **tabs**, (`\t`), a different separator can be specified with the `sep=","` option

```
datav <- read.table("bowens.csv",sep="," , header=T)
```

```
str(datav)
%> 'data.frame':      733 obs. of  3 variables:
%>  $ place: Factor w/ 727 levels "Abingdon","Admoor_Copse",...: 1 2 3 ..
%>  $ east : int  50 60 48 70 59 60 60 59 61 60 ...
%>  $ north: int  97 70 87 73 65 65 63 66 63 67 ...
```

read.table() : separators and decimal points

- the default field separator character in `read.table()` is `sep=" "`: which identifies with one or more spaces, one or more tabs (`\t`), and one or more newlines (`\n`)
- for comma-separated fields use `read.csv()`
- for semicolon-separated fields use `read.csv2()`
- for tab-delimited fields with decimal points as a commas, use `read.delim2()`

```
File: bowens.csv
```

```
-----  
|place,east,north |  
|Abingdon,50,97   |  
|Admoor Copse,60,70|  
|...             |  
|Youlbury,48,3    |  
-----
```

```
str(bw)  
%> 'data.frame':   733 obs. of  3 variables:  
%> $ place: Factor w/ 727 levels "AERE_Harwell",...: 2 3 1 4 5 ...  
%> $ east : int  50 60 48 70 59 60 60 59 61 60 ...  
%> $ north: int  97 70 87 73 65 65 63 66 63 67 ...
```

read.csv() and read.delim()

- additional functions to read a file in table format exist

```
> ?read.table  
...  
read.delim(file, header = TRUE, sep = "\t", quote = "\"",  
           dec = ".", fill = TRUE, comment.char = "", ...)  
...  
read.csv(file, header = TRUE, sep = ",", quote = "\"",  
         dec = ".", fill = TRUE, comment.char = "", ...)  
...  
read.csv2(file, header = TRUE, sep = ";", quote = "\"",  
          dec = ".", fill = TRUE, comment.char = "", ...)  
...  
read.delim2(file, header = TRUE, sep = "\t", quote = "\"",  
            dec = ",", fill = TRUE, comment.char = "", ...)
```

- further detailed instructions in the 'R Data Import/Export' manual:
<https://cran.r-project.org/doc/manuals/r-release/R-data.html>

- R can read data from the network using HTTP by specifying the file URL

```
wc <- read.table("https://tinyurl.com/murders-txt", header=T)

str(wc)
%> 'data.frame': 50 obs. of 4 variables:
%> $ state : Factor w/ 50 levels "Alabama","Alaska",...: 1 2 ...
%> $ population: int 3615 365 2212 2110 21198 2541 3100 ...
%> $ murder : num 15.1 11.3 7.8 10.1 10.3 6.8 3.1 6.2 ...
%> $ region : Factor w/ 4 levels "North.Central",...: 3 4 4 ...
```

- several packages available on CRAN to help R communicate with DBMSs: combining a unified 'front-end' package with a 'back-end' module, several common relational databases can be accessed (RMySQL, ROracle, RPostgreSQL and RSQLite)
- finally, R can read binary data files: NASA's HDF5 (Hierarchical Data Format, <https://www.hdfgroup.org/HDF5/>) and UCAR's netCDF data files (network Common Data Form, <http://www.unidata.ucar.edu/software/netcdf/>)
- and image files

Example: data Input from the Web

- let's retrieve the latest data on the COVID-19 Virus infection from the European Centers for Disease Control <https://www.ecdc.europa.eu/en>
- R can read data from the network using HTTP by specifying the file URL



European Centre for Disease Prevention and Control

An agency of the European Union

All topics: A to Z News & events Publications & data Tools About us Q

Coronavirus disease

The disease is rapidly spreading worldwide and the number of cases in Europe is rising with increasing pace in several affected areas.

[Latest information on COVID-19](#)

Coronavirus disease (COVID-19) COVID-19: Social distancing measures Antimicrobial resistance in zoonoses 2019/2020 influenza season

COVID-19

Several countries are now experiencing sustained local transmission of coronavirus disease 2019 (COVID-19), including Europe. The COVID-19 pandemic is rapidly evolving, and outbreak investigations are ongoing. ECDC is closely monitoring this outbreak, providing risk assessments, public health guidance, and advice on response activities to EU Member States and the EU Commission.

- [Latest situation update, epidemiological curve and global distribution](#)

Example: data Input from the Web

- we download an EXCEL file
- we use the following packages: `lubridate`, `curl` and `readxl`

```
url <- "https://www.ecdc.europa.eu/sites/default/files/documents/"
fname <- "COVID-19-geographic-disbtribution-worldwide-"
date <- lubridate::today() - 1
ext = ".xlsx"
target <- paste(url, fname, date, ext, sep="")
message("target:␣", target)

tmp_file <- tempfile("data", "/tmp", fileext=ext)
tmp <- curl::curl_download(target, destfile=tmp_file)
```

- data are imported in a tibble data structure

```
(data <- readxl::read_xlsx(tmp_file))
%> A tibble: 6,012 x 8
%>   DateRep          Day Month   Year Cases Deaths Countries. GeoId
%>   <dtm>          <dbl> <dbl> <dbl> <dbl>   <dbl> <chr>      <chr>
%> 1 2020-03-21 00:00:00   21     3  2020     2       0 Afghanistan AF
%> 2 2020-03-20 00:00:00   20     3  2020     0       0 Afghanistan AF
%> 3 2020-03-19 00:00:00   19     3  2020     0       0 Afghanistan AF
%> 4 2020-03-18 00:00:00   18     3  2020     1       0 Afghanistan AF
%> 5 2020-03-17 00:00:00   17     3  2020     5       0 Afghanistan AF
%> 6 2020-03-16 00:00:00   16     3  2020     6       0 Afghanistan AF
%> 7 2020-03-15 00:00:00   15     3  2020     3       0 Afghanistan AF
%> 8 2020-03-11 00:00:00   11     3  2020     3       0 Afghanistan AF
%> 9 2020-03-08 00:00:00    8     3  2020     3       0 Afghanistan AF
%>10 2020-03-02 00:00:00    2     3  2020     0       0 Afghanistan AF
% ... with 6,002 more rows
```

Homework

- learn how to extract and manipulate data from the imported tibble
- in the next lecture we will see how to represent data with histograms and plots