


Bayesian Networks

Alberto Garfagnini

Università di Padova

AA 2020/2021 - Stat Lect. 14



The Truck Driver Example

- a truck driver is due to make a 800 km trip
- we analyze the risk of his falling asleep while driving
- there may be **causal relationships** between:
 - 1) the **driver's sleep** (did he sleep well and for more than 7 hours the night before ?)
 - 2) his **perceived fatigue** (does he feel tired at the beginning of the trip ?)
 - 3) the **risk of falling asleep** while driving

Current situation

- the truck **driver feels tired** at the beginning of the trip
- weather or not this is due to a bad sleep the night before, or any other reason, is of no use to evaluate the risk
- the **driver feels perfectly fit** before starting to drive
- the quality of sleep the night before has no influence on his current condition
- the risk of falling asleep is **conditionally independent** of the quality of his sleep, given the driver's current fatigue

The Truck Driver Example

Our formal model:

- we model it with binary variables which tell us if
 - X_1 : the truck driver slept well the night before
 - X_2 : he feels tired at the beginning of the trip
 - X_3 : he will fall asleep while driving
- but, the quality of sleep the night before has no influence on his current condition

$$P(X_3 \mid X_1, X_2) = P(X_3 \mid X_2)$$

therefore:

$$\begin{aligned} P(X_1, X_2, X_3) &= P(X_1 \mid X_2, X_3) P(X_2, X_3) \\ &= P(X_3 \mid X_2) P(X_2 \mid X_1) P(X_1) \end{aligned}$$

The Doped Athlete Example

- during a sports competition, each athlete undergoes two doping tests
 - test **A is a blood test**
 - test **B a urine test**
- the two tests are carried out in two different laboratories, no contact between the to labs is possible

Current situation

- the results of the two tests are not independent variables:
 - if test A is positive → the participant is likely to have used a banned product → test B will probably be also positive
 - a participant has **has taken** and detectable substance
- tests A and B can be considered independent, since the two laboratories use different detection methods
 - a participant **has not taken** any prohibited substance
- tests A and B can be again considered independent → they may give a negative response according to the test efficacy
- the results of both tests are **conditionally independent**, given the status of the tested athlete

The Doped Athlete Example

Our formal model:

- we model it with binary variables which tell us if
 - X_1 : the athlete is *clean* or not
 - X_2 : the result of *test A*
 - X_3 : the result of *test B*
- let's write the conditional probability

$$P(X_3 \mid X_2, X_1) = P(X_3 \mid X_1)$$

- knowing whether the athlete has taken the substance is enough information to estimate the chances of test B being positive
- the symmetric equation holds

$$P(X_2 \mid X_3, X_1) = P(X_2 \mid X_1)$$

combining the results:

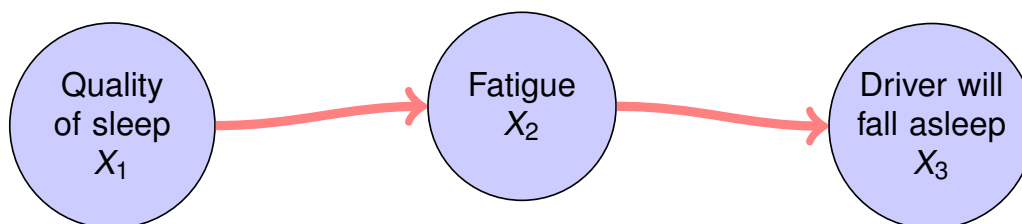
$$\begin{aligned} P(X_1, X_2, X_3) &= P(X_3 \mid X_2, X_1) P(X_2 \mid X_1) P(X_1) \\ &= P(X_3 \mid X_1) P(X_2 \mid X_1) P(X_1) \end{aligned}$$

Discrete Bayesian Networks

- both examples can be expressed in a Bayesian network

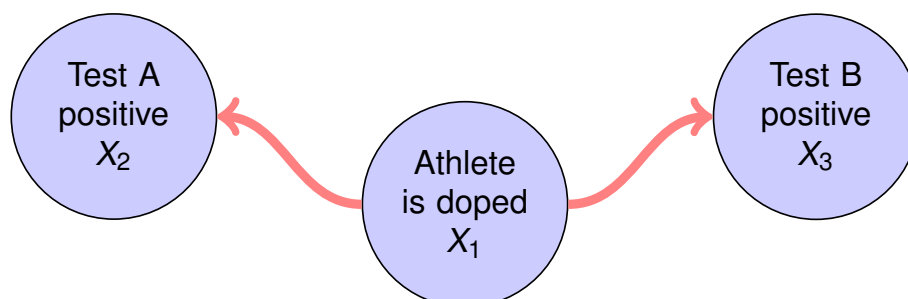
The Truck Driver Example

- that there is an influence of variable X_1 on variable X_2 , and of variable X_2 on variable X_3



The Doped Athlete Example

- X_2 and X_3 are conditionally independent given X_1



Bayesian Network


Definition

- given
 - n random variables X_1, X_2, \dots, X_n
 - a **directed acyclic graph with n numbered nodes**
 - suppose node j of the graph associated to the X_j variable
- the graph is a Bayesian network, representing the variables X_1, X_2, \dots, X_n , if

$$P(X_1, X_2, \dots, X_n) = \prod_{j=1}^n P(X_j \mid \text{parents}(X_j))$$

where $\text{parents}(X_j)$ denotes the set of all variables X_k , such that there is an arc from node k to node j

Proposition



- **any joint probability distribution may be represented by a Bayesian network**

$$\begin{aligned} P(X_1 \mid X_2, \dots, X_n) &= P(X_1) P(X_2 \dots X_n \mid X_1) \\ &= P(X_1) P(X_2 \dots X_1) P(X_3 \dots X_n \mid X_1 X_2) \\ &= \dots \\ &= P(X_1) P(X_2 \mid X_1) P(X_3 \mid X_2, X_1) \dots P(X_n \mid X_1, X_2, \dots, X_{n-1}) \end{aligned}$$

A. Garfagnini (UniPD)

AdvStat 4 PhysAna - Stat-Lec.14

6

Bayesian Networks in R: case study

The transportation means survey

- let's consider an hypothetical survey whose aim is to investigate the usage patterns of different means of transport, with a focus on private cars and public trains or buses
- each regular commuting individual fills a questionnaire on the following six discrete variables
 - Age (A)**: below 30 (young), between 30 and 60 (adult) greater than 60 (senior)
 - Sex (S)**: male (M) or female (F)
 - Education (E)**: highest individual degree between high school (high) and university or higher (uni)
 - Occupation (O)**: weather the individual is an employee (empl) or a self-employed worker (self)
 - Residence (R)**: the size of the city the individual lives in, a (small) or (big) town
 - Travel (T)**: the means of transport flavored by the individual, car, train or bus
- the variables can be grouped into **demographic indicators** (Age and Sex), **socioeconomic indicators** (Education, Occupation and Residence) and the **target of the survey** (Travel)

Travel Survey in bnlearn


- we use the bnlearn R package to build a directed acyclic graphs (DAG) that describes the network

```
library(bnlearn)
```

```
tus_dag <- empty.graph(nodes = c("A", "S",  
                                "E", "O", "R", "T"))  
tus_dag <- set.arc(tus_dag, from = "A", to = "E")  
tus_dag <- set.arc(tus_dag, from = "S", to = "E")  
tus_dag <- set.arc(tus_dag, from = "E", to = "O")  
tus_dag <- set.arc(tus_dag, from = "E", to = "R")  
tus_dag <- set.arc(tus_dag, from = "O", to = "T")  
tus_dag <- set.arc(tus_dag, from = "R", to = "T")
```

```
tus_dag  
#> Random/Generated Bayesian network  
#> model:  
#> [A][S][E|A:S][O|E][R|E][T|O:R]  
#> nodes:  
#> arcs:  
#> undirected arcs:  
#> directed arcs:  
#> average markov blanket size:  
#> average neighbourhood size:  
#> average branching factor:  
#> generation algorithm:
```

```
6  
6  
0  
6  
2.67  
2.00  
1.00  
Empty
```



Exploring the Travel Survey DAG bnlearn

- show the nodes of a graph:

```
nodes(tus_dag)  
[1] "A" "S" "E" "O" "R" "T"
```

- examine the nodes close to a target


```
# The neighbourhood of 'E'  
nbr(tus_dag, "E")  
#> [1] "A" "S" "O" "R"  
parents(tus_dag, "E")  
#> [1] "A" "S"  
children(tus_dag, "E")  
#> [1] "O" "R"
```

- look for roots (no parents) and leaves (no children) nodes:

```
root.nodes(tus_dag)  
#> [1] "A" "S"  
leaf.nodes(tus_dag)  
#> [1] "T"
```

- and plot the graph

```
library(Rgraphviz)  
graphviz.plot(plant)
```



Travel Survey in bnlearn

- another way to create the network is using the model formula interface provided by modelstring

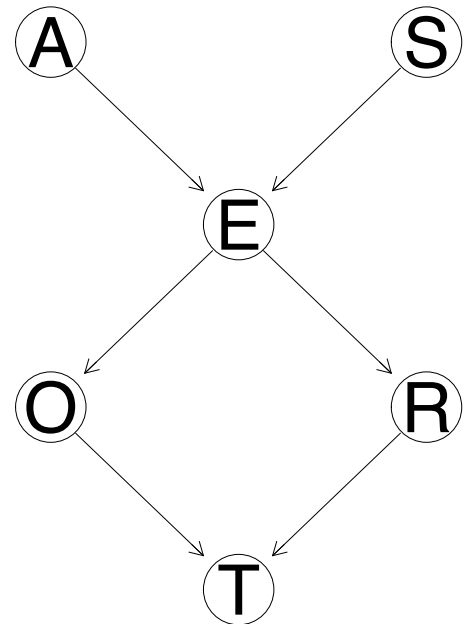
```
dag <- empty.graph(nodes = c("A", "S", "E", "O", "R", "T"))
```

```
dag2 <- model2network("[A][S][E|A:S][O|E][R|E][T|O:R]")
```

```
all.equal(tus_dag, dag2)
[1] TRUE
```

- we now define the **levels of the nodes**,
i.e. the discrete values defined on a non-ordered set

```
A_lvl1 <- c("young", "adult", "old")
S_lvl1 <- c("M", "F")
E_lvl1 <- c("high", "uni")
O_lvl1 <- c("emp", "self")
R_lvl1 <- c("small", "big")
T_lvl1 <- c("car", "train", "other")
```



Travel Survey in bnlearn

- To complete the BN modelling the survey → specify the **joint probabilities**

```
A_prob <- array(c(0.30, 0.50, 0.20), dim = 3,
               dimnames = list(A = A_lvl1))
S_prob <- array(c(0.60, 0.40), dim = 2, dimnames = list(S = S_lvl1))

O_prob <- array(c(0.96, 0.04, 0.92, 0.08), dim = c(2, 2),
               dimnames = list(O = O_lvl1, E = E_lvl1))
R_prob <- matrix(c(0.25, 0.75, 0.20, 0.80), ncol = 2,
                 dimnames = list(R = R_lvl1, E = E_lvl1))

E_prob <- array(c(0.75, 0.25, 0.72, 0.28, 0.88, 0.12,
                 0.64, 0.36, 0.70, 0.30, 0.90, 0.10), dim = c(2, 3, 2),
               dimnames = list(E = E_lvl1, A = A_lvl1, S = S_lvl1))
T_prob <- array(c(0.48, 0.42, 0.10, 0.56, 0.36, 0.08,
                 0.58, 0.24, 0.18, 0.70, 0.21, 0.09), dim = c(3, 2, 2),
               dimnames = list(T = T_lvl1, O = O_lvl1, R = R_lvl1))
```

- and finally **associate the probabilities** to the **BN model**

```
cpt <- list(A = A_prob, S = S_prob, E = E_prob, O = O_prob,
            R = R_prob, T = T_prob)
bn <- custom.fit(tus_dag, cpt)
```

- variables that are not linked by an arc are conditionally independent

→ we can factorise the global distribution

$$P(A, S, E, O, R, T) = P(A) \cdot P(S) \cdot P(E | A, S) \cdot P(O | E) \cdot P(R | E) \cdot P(T | O, R)$$

Travel Survey : estimate the probability table

- we knew the DAG and the probabilities defining the BN → BNs are used as expert systems
- but in most cases, the **parameters of the local distributions** will be **inferred** (i.e. learned) **from the observed sample**

```
survey <- read.table("survey.txt", header = TRUE,
                     stringsAsFactors = TRUE)
```

```
head(transp_survey, n=3)
#>      A      R      E      O S      T
#> 1 adult    big high emp F    car
#> 2 adult small uni emp M    car
#> 3 adult    big uni emp F train
```

```
tail(transp_survey, n=3)
#>      A      R      E      O S      T
#> 498   old big high emp M train
#> 499 adult big high emp F other
#> 500 adult big high emp M other
```

- the **conditional probabilities** can be estimated looking at the corresponding **empirical frequencies** in the data set

$$P(O = \text{emp} \mid E = \text{high}) = \frac{P(O = \text{emp}, E = \text{high})}{P(E = \text{high})}$$
$$= \frac{\text{number of observations for which } O = \text{emp and } E = \text{high}}{\text{number of observations for which } E = \text{high}}$$

Travel Survey : estimate the probability table

- the **bn.fit()** function computes the classic frequentist and maximum likelihood estimates from the data
- it complements the **custom.fit()** function which constructs a BN using a set of custom parameters specified by the user

```
travel_dag <- model2network("[A][S][E|A:S][O|E][R|E][T|O:R]")
```

```
bn.mle <- bn.fit(travel_dag, data = transp_survey,
                 method = "mle")
```

- as an alternative, the same conditional probabilities can be estimated in the Bayesian framework, using their posterior distributions

```
bn.bayes <- bn.fit(travel_dag, data = transp_survey,
                   method = "bayes", iss = 10)
```

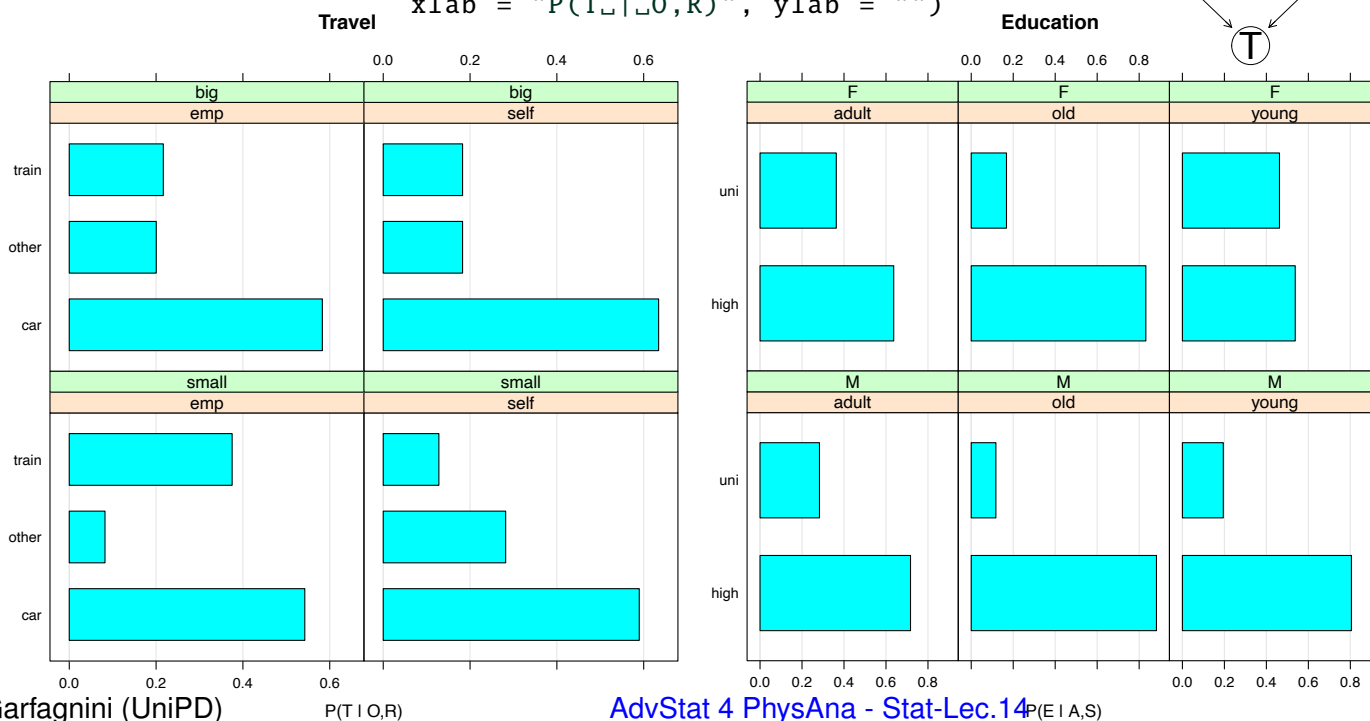
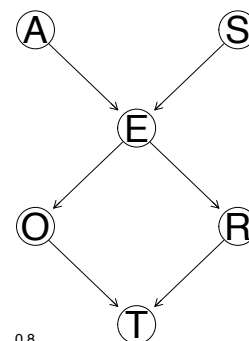
- the **iss** optional argument (imaginary sample size) determines how much weight is assigned to the prior distribution compared to the data when computing the posterior
- **iss** is typically chosen between 1 and 15, to allow the prior distribution to be easily dominated by the data (see bnlearn documentation)

Plotting the probability between links

- we can plot the probabilities associated to each link of the network

```
bn.fit.barchart(bn.bayes$E, main = "Education",
               xlab = "P(EU | A, S)", ylab = "")
```

```
bn.fit.barchart(bn.bayes$T, main = "Travel",
               xlab = "P(TU | O, R)", ylab = "")
```



A. Garfagnini (UniPD)

$P(T | O, R)$

AdvStat 4 PhysAna - Stat-Lec.14 $P(E | A, S)$

14

Investigating a DAG structure from data

- so far we have assumed that the **DAG underlying the BN is known**
- we rely on prior knowledge on the phenomenon to decide which arcs are present in the graph and which are not → **expert system**
- the structure of the DAG** itself may be the **object of our investigation**
- in genetics and systems biology it is common to reconstruct the molecular pathways and networks underlying complex diseases and metabolic processes
- learning the DAG of a BN is a complex task**
 - the space of the possible DAGs is very big → it grows exponentially with the number of nodes
 - this space is very different from real spaces : it is not continuous and has a finite number of elements → ad-hoc algorithms are required to explore it
- two classes of **statistical criteria used to evaluate DAGs**
 - conditional independence tests
 - network scores

Conditional Independence Test

- arcs encode a probabilistic dependence → conditional independence tests can be used to verify if that probabilistic dependence is supported by the data

E → T

- H_0 : Travel is independent of Education
- use the log-likelihood ratio G^2

$$P(T, E \mid O, R) = \sum_{t \in T} \sum_{e \in E} \sum_{k \in O \times R} n_{tek} \log \frac{n_{tek} n_{++k}}{n_{t+k} n_{+ek}},$$

- the use of a "+" subscript denotes the sum over that index :
- n_{t+k} = sum over the second index, $e \in E$
- n_{++k} = sum over the first ($t \in T$) and second ($e \in E$) indexes, respectively
- or Pearson's X^2 :

$$P(T, E \mid O, R) = \sum_{t \in T} \sum_{e \in E} \sum_{k \in O \times R} \frac{(n_{tek} - m_{tek})^2}{m_{tek}} \quad \text{with } m_{tek} = \frac{n_{t+k} n_{+ek}}{n_{++k}}$$

- both tests have an asymptotic χ^2 distribution under H_0 .

Travel Survey: evaluate DAG arcs

Mutual Information, log-likelihood ratio G^2

```
transp_survey <- read.table("survey.txt", header = TRUE,
                             stringsAsFactors = TRUE)

ci.test("T", "E", c("O", "R"), test = "mi", data = transp_survey)

#>      Mutual Information (disc.)
#>
#> data:  T ~ E | O + R
#> mi = 9.8836, df = 8, p-value = 0.2733
#> alternative hypothesis: true value is greater than 0
```

Pearson's X^2

```
ci.test("T", "E", c("O", "R"), test = "x2", data = transp_survey)

#>      Pearson's X^2
#>
#> data:  T ~ E | O + R
#> x2 = 8.2375, df = 8, p-value = 0.4106
#> alternative hypothesis: true value is greater than 0
```

- both tests return very large p-values → the dependence relationship encoded by $E \times T$ is not significant given the current DAG structure

- network scores focus on the DAG as a whole
- they provide a statistical measurement of **how well the DAG mirrors the dependence structure of the data**

Bayesian Information criterion (BIC)

$$\begin{aligned}\text{BIC} &= \log P(A, S, E, O, R, T) - \frac{d}{2} \log n \\ &= \log P(A) - \frac{d_A}{2} \log n + \log P(S) - \frac{d_S}{2} \log n \\ &\quad + \log P(E \mid A, S) - \frac{d_E}{2} \log n + \log P(O \mid E) - \frac{d_O}{2} \log n \\ &\quad + \log P(R \mid E) - \frac{d_R}{2} \log n + \log P(T \mid O, R) - \frac{d_T}{2} \log n\end{aligned}$$

- with n the sample size
- d the number of parameters of the network
- d_A, d_S, d_E, d_O, d_R and d_T the number of parameters associated with each node

Travel Survey: evaluate scores

several scores are available in bnlearn:

- Bayesian Information criterion (BIC)

```
score(travel_dag, data = transp_survey, type = "bic")  
#> [1] -2012.687
```

- Bayesian Dirichlet equivalent uniform (BDe)

```
score(travel_dag, data = transp_survey, type = "bde")  
#> [1] -2015.647
```

- using such scores it is possible to compare different DAGs and investigate which of them fits the data better

```
nparams(travel_dag, transp_survey)  
[1] 21
```

```
dag_new <- set.arc(travel_dag, from = "E", to = "T")  
nparams(dag_new, transp_survey)  
[1] 29
```

```
score(dag_new, data = transp_survey, type = "bic")  
[1] -2032.603
```

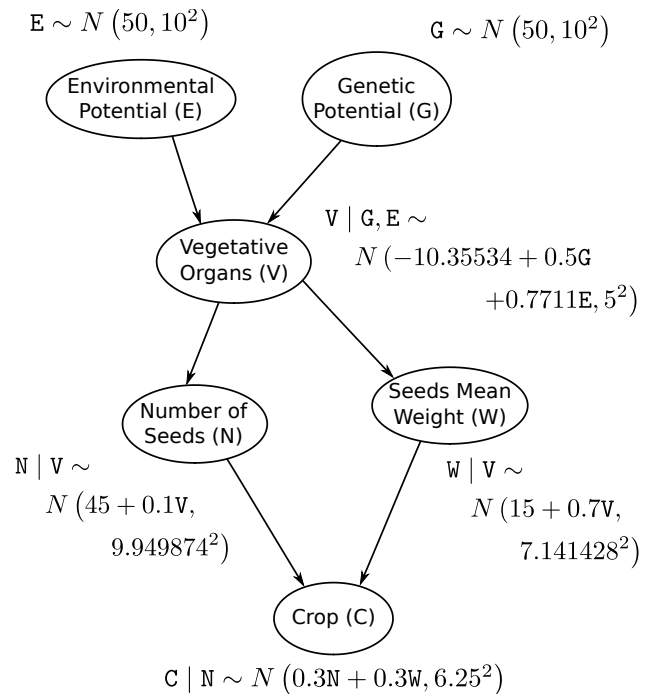
- adding **E → T** is not beneficial
the increase in $\log P(A, S, E, O, R, T)$ is not sufficient to offset the heavier penalty from the additional parameters

Continuous Bayesian Networks

- focus on modelling continuous data under a multivariate Normal distribution hypothesis
- we are interested in the analysis of a particular plant and study
 - the potential of the plant and the environment
 - the production of vegetative mass
 - the harvested grain mass, i.e. the *crop*

The plant network

- environmental potential (E)
- genetic potential (G)
- vegetative organs (V)
- number of seeds (N)
- seeds mean weight (W)
- crop (C)



A. Garfagnini (UniPD)

AdvStat 4 PhysAna - Stat-Lec.14

20

The Plant example: network structure

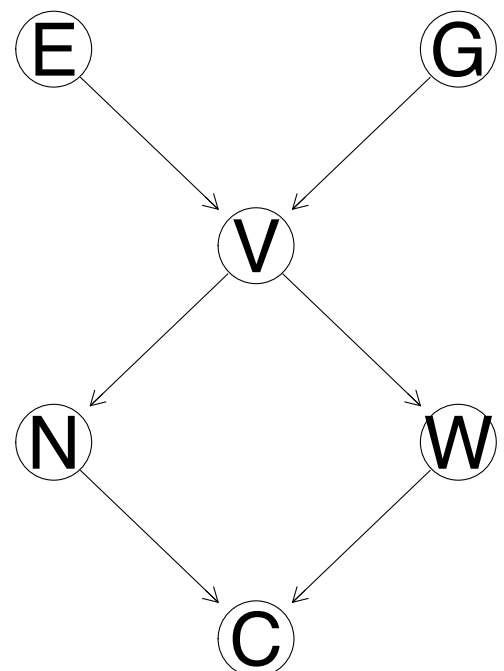
- let's build the DAG bayesian network:

```
plant_dag <- model2network("[G][E][V|G:E][N|V][W|V][C|N:W]")
plant_dag
```

```
#> Random/Generated Bayesian network

#> model:
#> [E][G][V|E:G][N|V][W|V][C|N:W]
#> nodes: 6
#> arcs: 6
#> undirected arcs: 0
#> directed arcs: 6
#> average markov blanket size: 2.67
#> average neighbourhood size: 2.00
#> average branching factor: 1.00
#> generation algorithm: Empty

nparams(plant)
#> [1] 18
```



A. Garfagnini (UniPD)

AdvStat 4 PhysAna - Stat-Lec.14

21

The Plant example: connection probabilities

- to make quantitative statements about the behavior of the variables in the BN, we need to completely specify their joint probability distribution
- if we are modelling n variables we must specify n means, n variances and $n(n-1)/2$ correlation coefficients

```
disE <- list(coef = c("(Intercept)" = 50), sd = 10)
disG <- list(coef = c("(Intercept)" = 50), sd = 10)
disV <- list(coef = c("(Intercept)" = -10.35534,
                     E = 0.70711, G = 0.5), sd = 5)
disN <- list(coef = c("(Intercept)" = 45,
                     V = 0.1), sd = 9.949874)
disW <- list(coef = c("(Intercept)" = 15,
                     V = 0.7), sd = 7.141428)
disC <- list(coef = c("(Intercept)" = 0,
                     N = 0.3, W = 0.7), sd = 6.25)
dis.list = list(E = disE, G = disG, V = disV,
               N = disN, W = disW, C = disC)

plant <- custom.fit(plant_dag, dist = dis.list)
```


References

R packages

- bnlearn
<https://cran.r-project.org/web/packages/bnlearn/>,
<https://www.bnlearn.com/>
- BayesNetBP
<https://cran.r-project.org/web/packages/BayesNetBP/index.html>
- bnstruct <https://github.com/sambofra/bnstruct>

Books

- O. Pourret et al, *Bayesian Networks, A Practical Guide to Applications*, J.Wiley and Sons, 2008, ISBN 978-0-470-06030-8
- M. Scutari, J.B. Denis, *Bayesian Networks with Examples in R*, CRC Press, 2015, ISBN 978-1-4822-2559-4



New R release 4.1.0

- released on May 18, 2021:

<https://cran.r-project.org/src/base/R-4/R-4.1.0.tar.gz>

From: Peter Dalgaard
Subject: [Rd] R 4.1.0 is released
Date: Tue May 18 10:06:24 CEST 2021
To: r-announce@r-project.org, Cc: r-devel@r-project.org

The build `system` rolled up `R-4.1.0.tar.gz` (codename "Camp_Pontanezen") this morning.

This `is` a major `update`, notably containing the `new` native pipe operator `"|>"` and shorthand inline functions `"\(\x)_x+1"`.

The `list` below details the changes in this release.

You can `get` the `source` code from

<https://cran.r-project.org/src/base/R-4/R-4.1.0.tar.gz>

`or` wait `for` it to be mirrored at a `CRAN` site nearer to you.

Binaries `for` various platforms will appear in due course.

`For` the `R` Core Team,
Peter Dalgaard

New R 4.1.0 features: the pipe operator, `|>`

- introduced with the `magrittr` package in 2014, and building upon Unix / linux scripting languages syntax and F# functional programming languages (like F#)
- decrease development time and improve code readability and maintainability
- Example:

- with `separate functions`:

```
raw_data = rnorm(100, mean = 4, sd = 1)
density_summary = density(raw_data); plot(density_summary)
```

- with `nested function calls`:

```
plot(
  density(
    rnorm(100, mean = 4, sd = 1)
  )
)
```

- with the `magrittr` package

```
# R version up to 4.0.5
library("magrittr")
rnorm(100, mean = 4, sd = 1) %>% density() %>% plot()
```

- with the new `R` pipe operator, `|>`:

```
# R version 4.1.0
rnorm(100, mean = 4, sd = 1) |> density() |> plot()
```

New R 4.1.0 features: %>% versus |>

- **Note:** |> is not a drop-in replacement for all uses of %>%
- The %>% allowed function calls to be written with or without parentheses

```
letters %>% head()
# [1] "a" "b" "c" "d" "e" "f"

letters %>% head
# [1] "a" "b" "c" "d" "e" "f"
```
- with the native pipe the parentheses must be present

```
# R 4.1.0: Make sure your parentheses are present:
letters |> head()
# [1] "a" "b" "c" "d" "e" "f"
```
- %>% allowed the caller to use the "piped-in" values anywhere in the function call, simply by using a dot (.) as a place holder

```
c("dogs", "cats", "rats") %>% grepl("at", .)
# [1] FALSE TRUE TRUE
```
- no place holder is provided with the native pipe, a function must be defined and used

```
find_at = function(x) grepl("at", x)
c("dogs", "cats", "rats") |> find_at()
# [1] FALSE TRUE TRUE

c("dogs", "cats", "rats") |>
  {function(x) grepl("at", x)}()
# [1] FALSE TRUE TRUE
```

New R 4.1.0 features: anonymous functions

- it is common practice to define anonymous function when using higher-order functions: in the purrr package (map(), reduce(), keep()) and in base-R: Map(), lapply()

```
# For each letter, find the name of each dataset
# in the {datasets} package that starts with that letter
purrr::map(
  letters[1:3],
  function(x) {
    ds = ls("package:datasets")
    ds[stringr::str_starts(tolower(ds), x)]
  }
)
#[[1]]
#[1] "ability.cov" "airmiles" "AirPassengers" "airquality"
...
#[[3]]
#[1] "cars" "ChickWeight" "chickwts" "co2" "CO2" "crimtab"
```
- or in base-R:

```
# In base R
Map(
  function(x) {
    pattern = paste0("^", x) # eg "^a" to match a leading 'a'
    grep(pattern, ls("package:datasets"), value = TRUE, ignore.case = TRUE)
  },
  letters[1:3]
)
```

New R 4.1.0 features: anonymous functions

- a new syntax has been introduced into R that may make these anonymous function declarations more succinct:

```
# R 4.1.0
Map(
  \(x) {
    pattern = paste0("^", x)
    grep(pattern, ls("package:datasets"), value = TRUE, ignore.case = TRUE)
  },
  letters[1:3]
)
```

- this syntax can help in accessing single values with the native pipe:

```
c("dogs", "cats", "rats") |> {\(x) grep1("at", x)}()
```

- the `\(x)` syntax can be used to write functions anywhere is needed

```
my_func <- \(x) mean(x, na.rm = TRUE)
my_func(c(1, 4, 7))
```

- and multiple arguments are allowed:

```
three_args <- \(a, b, c) a * b / c
three_args(1.5, 4, 3)
```

New R 4.1.0: concatenating factors

- in earlier R versions, combining factors with the syntax `c(factor1, factor2)` brings to a vector of integers

```
# with different levels:
factor("a")
# [1] a
# Levels: a
factor("b")
# [1] b
# Levels: b

c(factor("a"), factor("b"))
# [1] 1 1
```

```
# with identical levels:
c(factor("a"), factor("a"))
# [1] 1 1
```

- this happens because factors are stored internally as integers
- in $R \geq 4.1.0$, combining two factors together generates a new factor with levels that are a combination of the levels in the original factors

```
fac1 = factor(c("a", "b", "d"))
fac2 = factor(c("b", "c"))
c(fac1, fac2)

# [1] a b d b c
# Levels: a b d c
```